

CGS698C, Assignment 3

Jiyanshu Dhaka 220481

Part 1: Estimating the Posterior Distribution Using Different Computational Methods

```
[107]: import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
from scipy.special import factorial
from scipy.stats import norm
import seaborn as sns

exp = [10, 15, 15, 14, 14, 14, 13, 11, 12, 16]
s = 10
N = 20

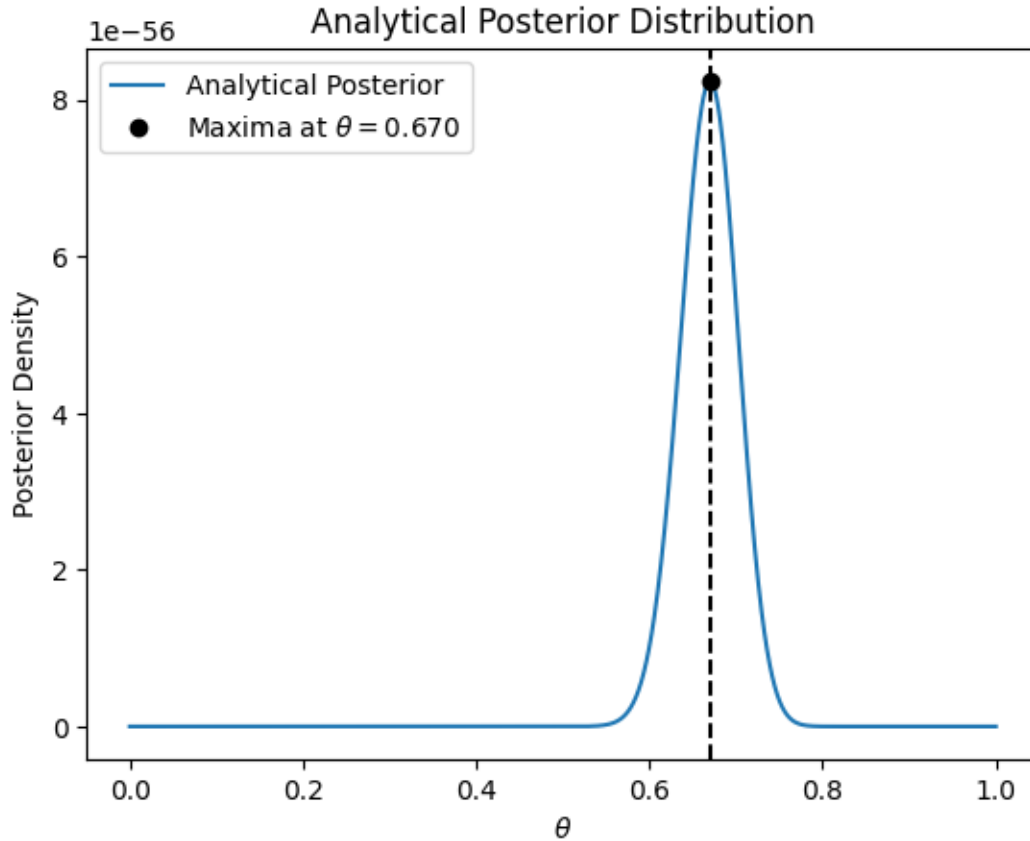
def pe(theta, alpha, beta):
    return (theta ** (alpha - 1)) * ((1 - theta) ** (beta - 1))

def le(theta, N, e):
    return ((theta ** e) * ((1 - theta) ** (N - e))) * (factorial(N) /
    ↪ (factorial(e) * factorial(N - e)))
```

1.1 Graph Analytical Posterior

```
[108]: theta_s = np.linspace(0,1, 1000)
y_pa = [pe(theta, 135, 67) for theta in theta_s]
max_i = np.argmax(y_pa)
theta_max = theta_s[max_i]
pdf_max = y_pa[max_i]
plt.plot(theta_s, y_pa, label="Analytical Posterior")
plt.plot(theta_max, pdf_max, 'ok', label=f'Maxima at  $\theta = \{theta\_max:.3f\}$ ')
plt.axvline(x=theta_max, color='k', linestyle='--')
plt.xlabel(" $\theta$ ")
```

```
plt.ylabel("Posterior Density")
plt.title("Analytical Posterior Distribution")
plt.legend()
plt.show()
```



1.2 Estimate Posterior with Grid


```
[109]: lb = 0.
ub = 1.
samples_ = 100
grid_s = [lb + ((ub - lb)*i)/samples_ for i in range(samples_)]
y_l = np.ones(samples_)
fig, ax = plt.subplots(1, 2, figsize=(12, 5))
for e in exp:
    for i in range(samples_):
        y_l[i] *= le(grid_s[i], N, e)
    y_post_i = y_l/(np.sum(y_l))
    ax[0].plot(grid_s, y_post_i, color = 'b', label = None)
ax[0].set_xlabel("$\\theta$")
```

```

ax[0].set_ylabel("Posterior Density")
ax[0].set_title("(Intermediate Grid Steps) Posterior Estimation")

Approx_ML = np.sum(y_l*1)
y_post_grid = y_l/Approx_ML
max_i = np.argmax(y_post_grid)
theta_max = grid_s[max_i]
pdf_max = y_post_grid[max_i]
ax[1].plot(grid_s, y_post_grid, color = 'r', linestyle = 'solid', linewidth = 3, label="Posterior")
ax[1].plot(theta_max, pdf_max, 'ok', label=f'Maxima at  $\theta = \{theta\_max:.3f\}$ ')
ax[1].axvline(x=theta_max, color='k', linestyle='--')
ax[1].set_xlabel(" $\theta$ ")
ax[1].set_ylabel("Posterior Density")
ax[1].set_title("Grid Approximation Method")
ax[1].legend()
plt.tight_layout()
plt.show()

```



1.3 Estimate Marginal Likelihood

```

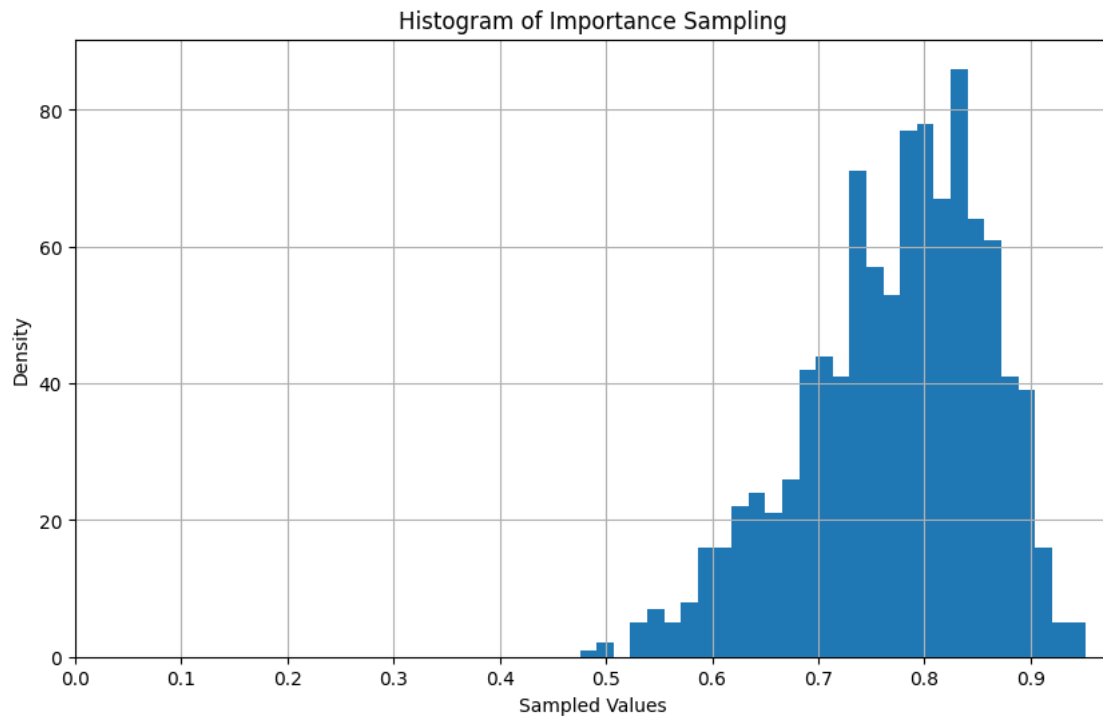
[110]: theta_s = np.random.beta(1, 1, 100000)
y_l = np.ones(100000)
for e in exp:
    for i in range(100000):
        y_l[i] *= le(theta_s[i], N, e)
ML = np.mean(y_l)
print("Average Likelihood = ", ML)

```

Average Likelihood = 1.423430051730923e-10

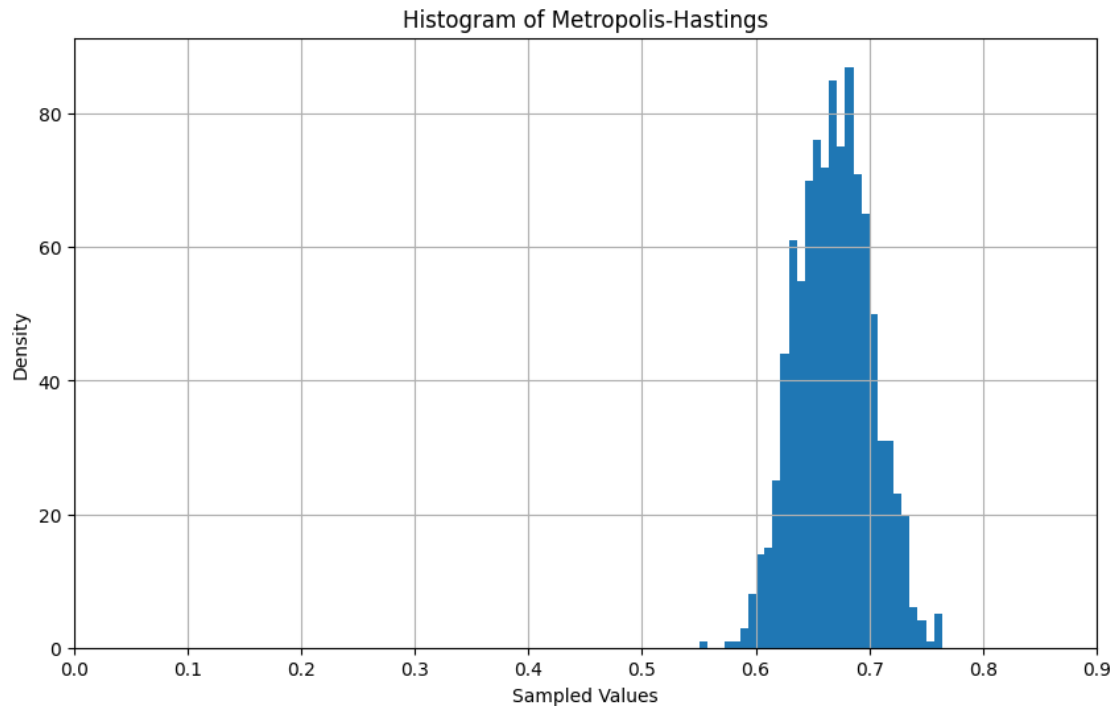
1.4 Importance Sampling

```
[111]: n = 4000
theta_s= abs(np.random.normal(0.,1., n))
DF = pd.DataFrame(columns = ['theta', 'weight'], index= range(n))
for e in exp:
    for i in range(n):
        DF.loc[i, 'theta'] = theta_s[i]
        DF.loc[i, 'weight'] = (1e(theta_s[i], N, e)*(theta_s[i] <= 1.))/norm.
        pdf(theta_s[i], 0., 1.)
y_posterior_imp = np.random.choice(DF['theta'],size = n//4, p = DF['weight']/np.
    sum(DF['weight']))
plt.figure(figsize=(10, 6))
plt.hist(y_posterior_imp, bins=30, label = 'density')
plt.title("Histogram of Importance Sampling")
plt.xlabel("Sampled Values")
plt.ylabel("Density")
plt.grid(True)
plt.xticks(np.arange(0., 1., 0.1))
plt.show()
```



1.5 Markov Chain Monte Carlo

```
[112]: num_samples = 1000
theta_s = np.zeros(num_samples)
theta_s[0] = np.random.beta(1, 1)
prev_post = 0.
for e in exp:
    prev_post = prev_post + np.log(le(theta_s[0], N, e))
i = 1
step = 0.04
reject = 0
while i < num_samples:
    proposed_theta = np.random.normal(theta_s[i-1], step)
    if(0. <= proposed_theta <= 1.):
        proposed_posterior = 0.
        for e in exp:
            proposed_posterior = proposed_posterior + np.log(le(proposed_theta,
↪N , e))
        h_r = proposed_posterior - prev_post
        hastings_ratio = np.exp(h_r)
        p_str = min(1, hastings_ratio)
        if(np.random.uniform(0, 1, 1) < p_str):
            theta_s[i] = proposed_theta
            prev_post = proposed_posterior
            i += 1
        else:
            reject += 1
    else:
        reject += 1
plt.figure(figsize=(10, 6))
plt.hist(theta_s, bins=30, label = 'density')
plt.title("Histogram of Metropolis-Hastings")
plt.xlabel("Sampled Values")
plt.ylabel("Density")
plt.grid(True)
plt.xticks(np.arange(0., 1., 0.1))
plt.show()
print("Rejections rate= ", (reject/num_samples)*100)
y_posterior_mc = theta_s
```

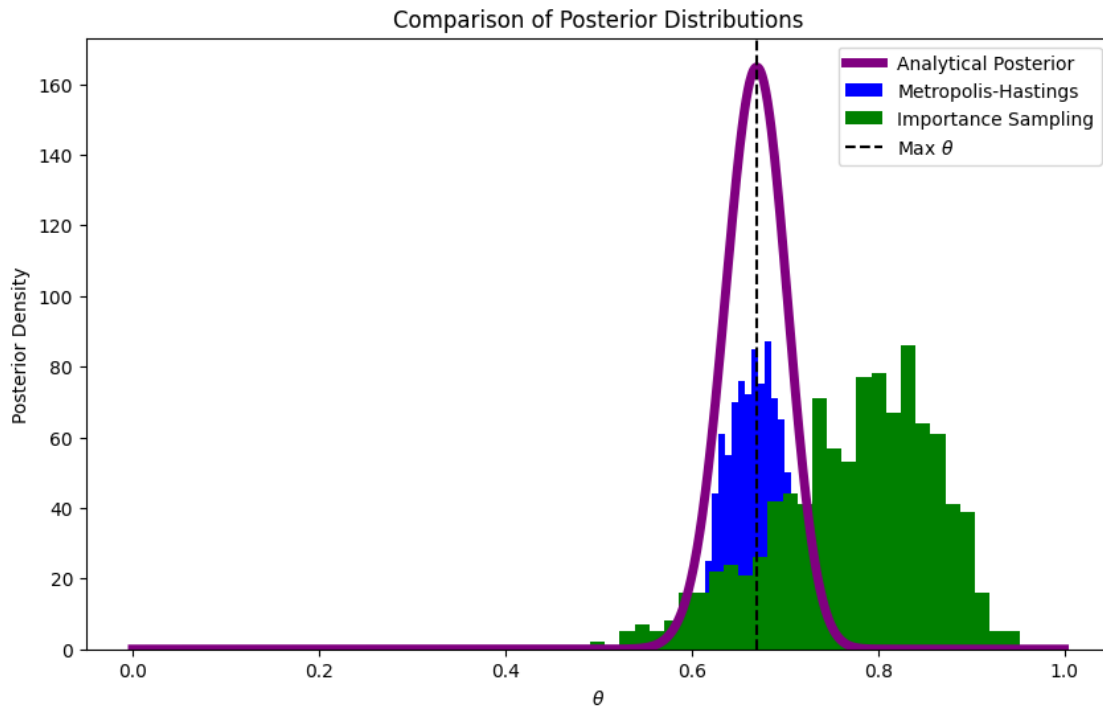


Rejections rate= 49.0

1.6 Comparison of Posterior Distributions

```
[113]: theta_s = np.linspace(0, 1, 1000)
y_pa_scaled = [pe(theta, 135, 67)*(2*(1e+57)) for theta in theta_s]
plt.figure(figsize=(10, 6))
plt.plot(theta_s, y_pa_scaled, label="Analytical Posterior", linewidth=5,
         color='purple')
plt.hist(y_posterior_mc, color='blue', bins=30, label="Metropolis-Hastings")
plt.hist(y_posterior_imp, color='green', bins=30, label="Importance Sampling")
plt.axvline(x=theta_max, color='k', linestyle='--', label="Max  $\theta$ ")
plt.xlabel(" $\theta$ ")
plt.ylabel("Posterior Density")
plt.title("Comparison of Posterior Distributions")
plt.legend()
plt.show()
print("Here, I have scaled Analytical Posterior for better analysis of the
      graphs.")
print("From the graph:")
print("(1). Metropolis-Hastings Method fits better to the Analytical Method.")
print("(2). Maxima for Metropolis-Hastings Method and Analytical Method are
      nearly the same.")
```

```
print("(3). Importance Sampling fits less to both of the other methods.")
```



Here, I have scaled Analytical Posterior for better analysis of the graphs.
From the graph:

- (1). Metropolis-Hastings Method fits better to the Analytical Method.
- (2). Maxima for Metropolis-Hastings Method and Analytical Method are nearly the same.
- (3). Importance Sampling fits less to both of the other methods.

Part 2: Writing your own sampler for Bayesian inference

2.5.1 Markov Chain Monte Carlo

```
[114]: import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
from scipy.stats import norm

url = "https://raw.githubusercontent.com/yadavhimanshu059/CGS698C/main/notes/
↳Data/word-recognition-times.csv"
data = pd.read_csv(url)

n = 4000
```

```

s = 30.

def posterior(data, a, b):
    t = (data['type'] == 'non-word').astype(float)
    m = a + t*b
    p = norm.logpdf(data['RT'], m, s) + norm.logpdf(a, 400, 50) + (t)*abs(norm.
↪logpdf(b, 0, 50))
    p = np.sum(p)
    return p

i = 1
step = 0.08
reject = 0
num_samples = 10000
a_samples = np.zeros(num_samples)
b_samples = np.zeros(num_samples)
a_samples[0] = np.random.normal(400, 50)
b_samples[0] = abs(np.random.normal(0, 50))
prev_post = posterior(data, a_samples[0], b_samples[0])

while i < num_samples:
    proposed_a = np.random.normal(a_samples[i-1], step)
    proposed_b = abs(np.random.normal(b_samples[i-1], step))
    post_new = posterior(data, proposed_a, proposed_b)
    h_r = post_new - prev_post
    hastings_ratio = np.exp(h_r)
    p_str = min(1, hastings_ratio)

    if(np.random.uniform(0, 1) < p_str):
        a_samples[i] = proposed_a
        b_samples[i] = proposed_b
        prev_post = post_new
        i += 1
    else:
        reject += 1

print(f"Rejections rate= {(reject/n)*100}")

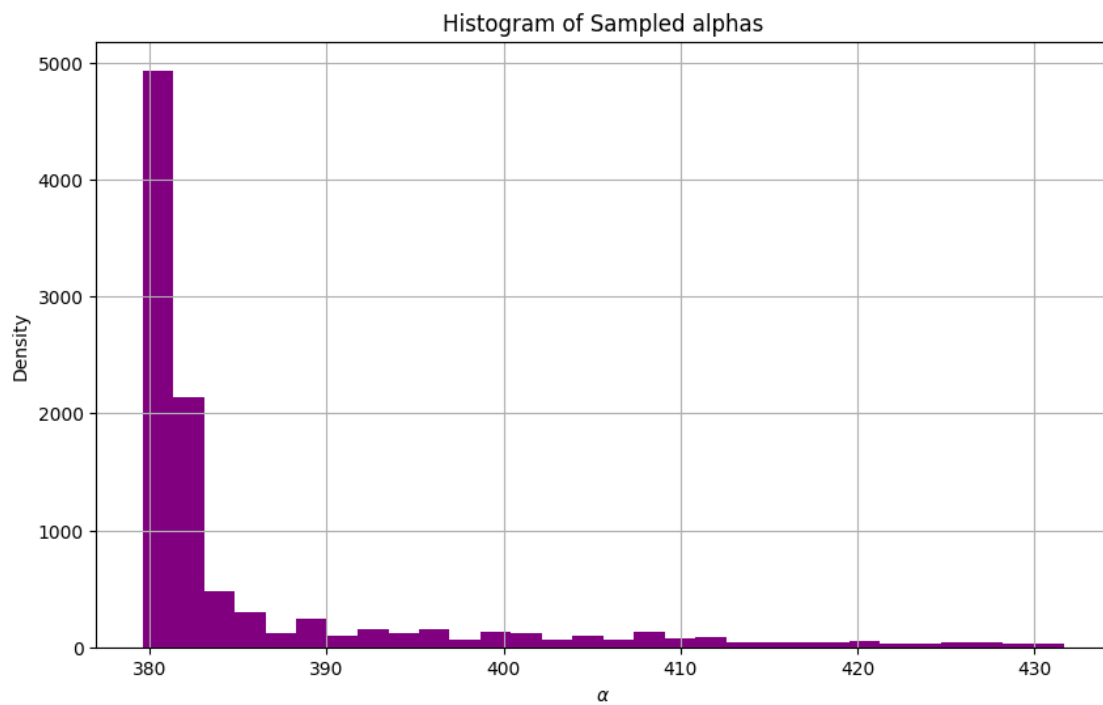
```

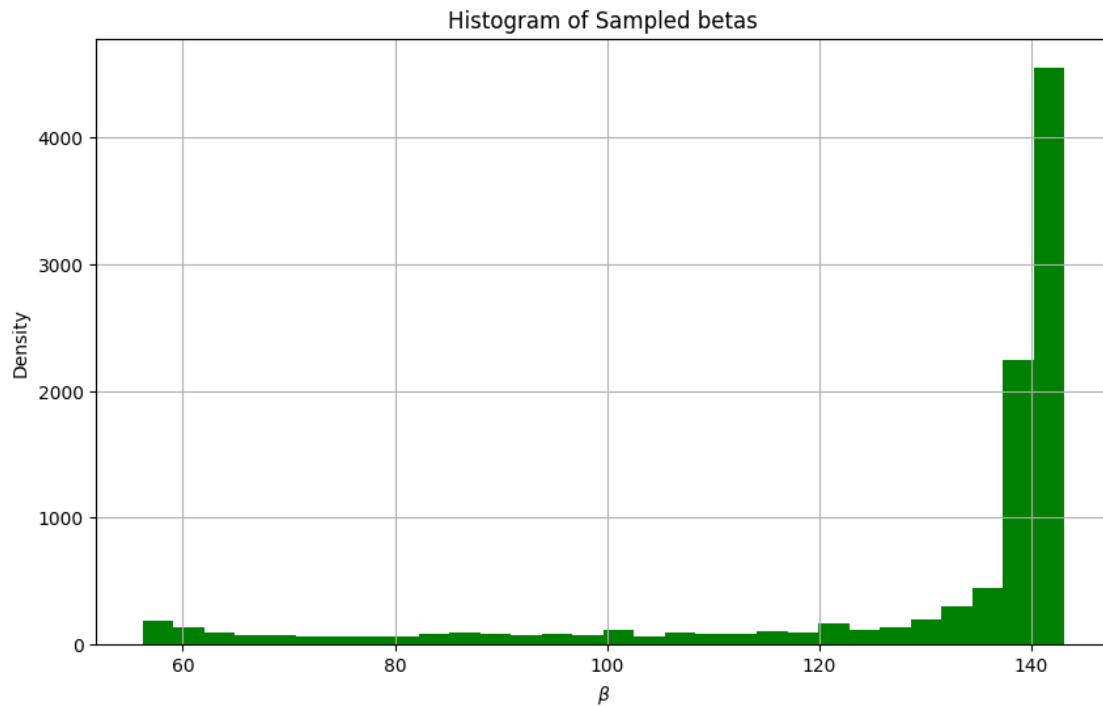
Rejections rate= 44.574999999999996

2.5.2 Histogram Plotting

```
[115]: plt.figure(figsize=(10, 6))
plt.hist(a_samples, bins=30, color='purple', label='alpha')
plt.title("Histogram of Sampled alphas")
plt.xlabel("$\\alpha$")
plt.ylabel("Density")
plt.grid(True)
plt.show()

plt.figure(figsize=(10, 6))
plt.hist(b_samples, bins=30, color='green', label='beta')
plt.title("Histogram of Sampled betas")
plt.xlabel("$\\beta$")
plt.ylabel("Density")
plt.grid(True)
plt.show()
```





Credible Intervals

```
[116]: burn = 1000
a_samples = a_samples[burn:]
b_samples = b_samples[burn:]
a_cred_interval = np.percentile(a_samples, [2.5, 97.5])
b_cred_interval = np.percentile(b_samples, [2.5, 97.5])

print(f"95% credible interval for alpha: {a_cred_interval}")
print(f"95% credible interval for beta: {b_cred_interval}")
```

95% credible interval for alpha: [380.20700143 397.71151624]

95% credible interval for beta: [98.9429619 142.26249731]

Part 3: Hamiltonian Monte Carlo sampler

3.1 HMC sampler

```
[117]: mu_true = 800
sigma_true = 100
data = np.random.normal(mu_true, np.sqrt(sigma_true), 500)
```

```

def grad(mu, sigma, y, n, m, s, a, b):
    grad_mu = ((n*mu) - np.sum(y))/(sigma**2) + ((mu - m)/(s**2))
    grad_sigma = (n/sigma) - (np.sum((y - mu)**2)/(sigma**3)) + ((sigma - a)/
↪(b**2))
    return np.array([grad_mu, grad_sigma])

def V(mu, sigma, y, n, m, s, a, b):
    nlpd = -(np.sum(norm.logpdf(y, mu, sigma)) + norm.logpdf(mu, m, s) + norm.
↪logpdf(sigma, a, b))
    return nlpd

def HMC(data, n, m, s, a, b, step, L, initial_q, nsamp, nburn):
    mu_samples = np.empty(nsamp)
    sigma_samples = np.empty(nsamp)
    reject = 0
    mu_samples[0] = initial_q[0]
    sigma_samples[0] = initial_q[1]
    i = 0
    while i < nsamp - 1:
        q = np.array([mu_samples[i], sigma_samples[i]])
        p = np.random.normal(0, 1, size=len(q))
        current_q = q.copy()
        current_p = p.copy()
        current_V = V(current_q[0], current_q[1], data, n, m, s, a, b)
        current_T = np.sum(current_p**2)/2
        for l in range(L):
            p -= ((step/2)*grad(q[0], q[1], data, n, m, s, a, b))
            q += (step*p)
            p -= ((step/2)*grad(q[0], q[1], data, n, m, s, a, b))
        proposed_q = q
        proposed_p = p
        proposed_V = V(proposed_q[0], proposed_q[1], data, n, m, s, a, b)
        proposed_T = np.sum(proposed_p**2)/2
        delta_energy = (current_V + current_T) - (proposed_V + proposed_T)
        if(delta_energy < 0.):
            mu_samples[i+1] = proposed_q[0]
            sigma_samples[i+1] = proposed_q[1]
            i += 1
        else:
            if(delta_energy > 100):
                accept_prob = 1
            else:
                accept_prob = min(1, np.exp(delta_energy))
            if(accept_prob > np.random.uniform(0,1)):
                mu_samples[i+1] = proposed_q[0]
                sigma_samples[i+1] = proposed_q[1]
                i += 1

```

```

        else:
            reject += 1
            posteriors = pd.DataFrame({'mu_samples': mu_samples[nburn:],
            ↪ 'sigma_samples': sigma_samples[nburn:]})
            posteriors['sample_id'] = np.arange(1, len(posterior) + 1)
            return posteriors

n = len(data)
mean_prior = 1000
std_prior = 100
a_prior = 10
b_prior = 2
step_size = 0.02
leapfrog_steps = 12
initial_values = [1000, 11]
burn_samples = 2000

posteriors = HMC(data, n, mean_prior, std_prior, a_prior, b_prior, step_size,
    ↪ leapfrog_steps, initial_values, 6000, burn_samples)


mean_mu_posterior = np.mean(posterior['mu_samples'])
mean_sigma_posterior = np.mean(posterior['sigma_samples'])

fig, axes = plt.subplots(1, 2, figsize=(10, 5))
axes[0].hist(posterior['mu_samples'], bins=30, label='$\\mu$')
axes[0].axvline(x=mean_mu_posterior, color='r', linestyle='--', label=f'Mean =
    ↪ {mean_mu_posterior}')
axes[0].set_xlabel('$\\mu$')
axes[0].set_ylabel('Density')
axes[0].set_title('Samples of $\\mu$')
axes[0].legend()

axes[1].hist(posterior['sigma_samples'], bins=30, label='$\\sigma$')
axes[1].axvline(x=mean_sigma_posterior, color='r', linestyle='--', label=f'Mean
    ↪ {mean_sigma_posterior}')
axes[1].set_xlabel('$\\sigma$')
axes[1].set_ylabel('Density')
axes[1].set_title('Samples of $\\sigma$')
axes[1].legend()

plt.tight_layout()
plt.show()
print("\nConclusions for Exercise 3.1:")
print("The posterior distributions for both  $\mu$  and  $\sigma$  have converged to certain
    ↪ values.")
print("The mean values of the posteriors provide estimates for the parameters  $\mu$ 
    ↪ and  $\sigma$  .")

```



Conclusions for Exercise 3.1:

The posterior distributions for both μ and σ have converged to certain values. The mean values of the posteriors provide estimates for the parameters μ and σ .

3.2 posterior sensitivity

```
[130]: n_ = [100, 1000, 6000]
for n in n_:
    posteriors_ = HMC(data, n, mean_prior, std_prior, a_prior, b_prior,
    ↪ step_size, leapfrog_steps, initial_values, n, n // 3)
    print(f"Posterior Distribution for n = {n}\n")

    mean_mu = np.mean(posteriors_['mu_samples'])
    mean_sigma = np.mean(posteriors_['sigma_samples'])

    fig, axes = plt.subplots(1, 2, figsize=(10, 5))
    axes[0].hist(posteriors_['mu_samples'], bins=30, label='$\mu$')
    axes[0].axvline(x=mean_mu, color='r', linestyle='--', label=f'Mean = {
    ↪ mean_mu}')
    axes[0].set_xlabel('$\mu$')
    axes[0].set_ylabel('Density')
    axes[0].set_title(f'Samples of $\mu$ for n = {n}')
    axes[0].legend()

    axes[1].hist(posteriors_['sigma_samples'], bins=30, label='$\sigma$')
    axes[1].axvline(x=mean_sigma, color='r', linestyle='--', label=f'Mean = {
    ↪ mean_sigma}')
```


```

axes[1].set_xlabel('$\\sigma$')
axes[1].set_ylabel('Density')
axes[1].set_title(f'Samples of $\\sigma$ for n = {n}')
axes[1].legend()


plt.tight_layout()
plt.show()

```


Posterior Distribution for n = 100



Posterior Distribution for n = 1000



Posterior Distribution for n = 6000



```
[131]: print("\nHere are the findings from the graphs:")
        print("(1). For Large nsamp values (6000), the posterior is more precise_
        ↪and less steep.")
        print("(2). For nsamp values around the same order (1000, 6000), the mean_
        ↪value remains close.")
        print("(3). For smaller nsamp values (100), the values are more random due_
        ↪to fewer samples after burn-in. Drawing conclusions on mean or sd values is_
        ↪unreliable with very small samples.")
        print("Conclusion: As the sample size increases, the posterior_
        ↪distributions become more stable and the estimates for μ and σ tend to_
        ↪converge.")
```

Here are the findings from the graphs:

- (1). For Large nsamp values (6000), the posterior is more precise and less steep.
 - (2). For nsamp values around the same order (1000, 6000), the mean value remains close.
 - (3). For smaller nsamp values (100), the values are more random due to fewer samples after burn-in. Drawing conclusions on mean or sd values is unreliable with very small samples.
- Conclusion: As the sample size increases, the posterior distributions become more stable and the estimates for μ and σ tend to converge.

3.3 posteriors change with change in step-size parameter

```
[121]: step_ = [0.001, 0.005, 0.02]
for step in step_:
    posteriors_ = HMC(data, n, mean_prior, std_prior, a_prior, b_prior, step,
    ↪leapfrog_steps, initial_values, 6000, burn_samples)
    print(f"Posterior Distribution for step = {step}\n")


    mean_mu = np.mean(posteriors_['mu_samples'])
    mean_sigma = np.mean(posteriors_['sigma_samples'])

    fig, axes = plt.subplots(1, 2, figsize=(10, 5))
    axes[0].hist(posteriors_['mu_samples'], bins=30, label='$\mu$')
    axes[0].axvline(x=mean_mu, color='r', linestyle='--', label=f'Mean =
    ↪{mean_mu}')
    axes[0].set_xlabel('$\mu$')
    axes[0].set_ylabel('Density')
    axes[0].set_title(f'Samples of $\mu$ for step = {step}')
    axes[0].legend()


    axes[1].hist(posteriors_['sigma_samples'], bins=30, label='$\sigma$')
    axes[1].axvline(x=mean_sigma, color='r', linestyle='--', label=f'Mean =
    ↪{mean_sigma}')
    axes[1].set_xlabel('$\sigma$')
    axes[1].set_ylabel('Density')
    axes[1].set_title(f'Samples of $\sigma$ for step = {step}')
    axes[1].legend()

    plt.tight_layout()
    plt.show()
```


Posterior Distribution for step = 0.001



Posterior Distribution for step = 0.005



Posterior Distribution for step = 0.02



```
[125]: print("\nBy comparing step sizes:")
print("(1). A very small step-size leads to sampling similar to smaller nsamp.")
print("(2). A significantly large step-size (0.2) results in sampling similar_
↳to larger nsamp.")
print("(3). Increasing step-size shifts the mean of the distribution left for_
↳mu_samples, but not as much for sigma_samples.")
print("(4). For step sizes 0.001 and 0.005, samples follow a linear path, while_
↳for 0.2, samples traverse a wider range. This is evident in the 'Trace Plot'.
↳")
print("Conclusion: Optimal step sizes balance exploration and exploitation of_
↳the parameter space in HMC sampling.")
```

By comparing step sizes:

- (1). A very small step-size leads to sampling similar to smaller nsamp.
 - (2). A significantly large step-size (0.2) results in sampling similar to larger nsamp.
 - (3). Increasing step-size shifts the mean of the distribution left for mu_samples, but not as much for sigma_samples.
 - (4). For step sizes 0.001 and 0.005, samples follow a linear path, while for 0.2, samples traverse a wider range. This is evident in the 'Trace Plot'.
- Conclusion: Optimal step sizes balance exploration and exploitation of the parameter space in HMC sampling.

3.4 Visually inspect the mu and sigma chains


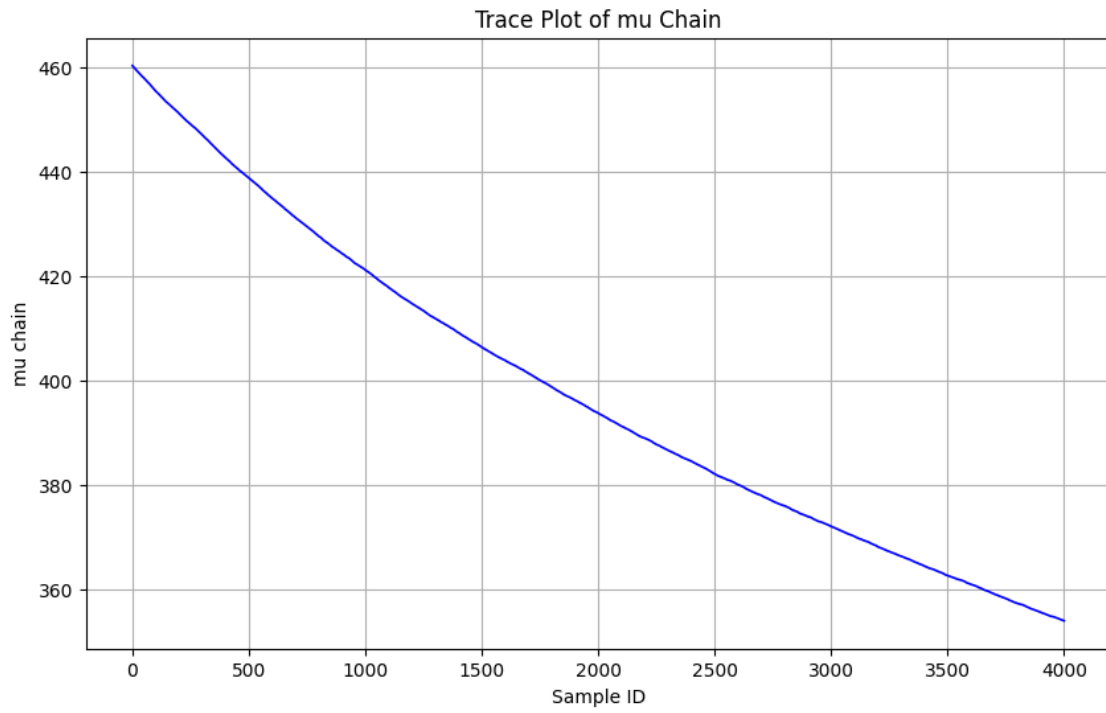
```
[126]: import seaborn as sns

step_sizes = [0.001, 0.005, 0.02]
for step in step_sizes:
    posteriors_ = HMC(data, n, mean_prior, std_prior, a_prior, b_prior, step,
    ↪leapfrog_steps, initial_values, 6000, burn_samples)
    print(f"Posterior Distribution for step = {step}\n")


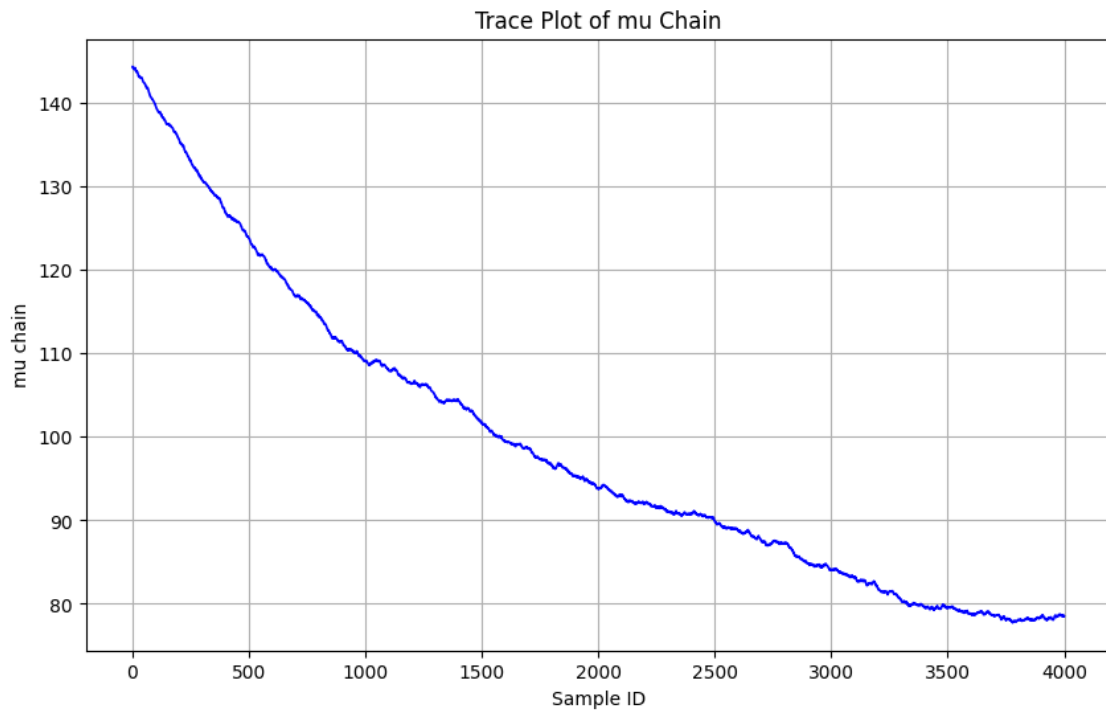
    plt.figure(figsize=(10, 6))
    sns.lineplot(data=posteriors_, x='sample_id', y='mu_samples', color='blue',
    ↪linewidth=1.2)
    plt.xlabel("Sample ID")
    plt.ylabel("mu chain")
    plt.title("Trace Plot of mu Chain")
    plt.grid(True)
    plt.show()

    plt.figure(figsize=(10, 6))
    sns.lineplot(data=posteriors_, x='sample_id', y='sigma_samples',
    ↪color='blue', linewidth=1.2)
    plt.xlabel("Sample ID")
    plt.ylabel("Sigma chain")
    plt.title("Trace Plot of sigma Chain")
    plt.grid(True)
    plt.show()
```



Posterior Distribution for step = 0.001



Posterior Distribution for step = 0.005



Posterior Distribution for step = 0.02



```
[127]: print("- Increasing step_size leads to less steep mu_samples and a leftward_
↳ shift in their mean. But for sigma_samples, the mean increases from 0.001 to_
↳ 0.005.")
print("Posterior Distribution for step = 0.001")
print("Posterior Distribution for step = 0.005")
print("Posterior Distribution for step = 0.02")
```

- Increasing step_size leads to less steep mu_samples and a leftward shift in their mean. But for sigma_samples, the mean increases from 0.001 to 0.005.

Posterior Distribution for step = 0.001

Posterior Distribution for step = 0.005

Posterior Distribution for step = 0.02

3.5 prior sensitivity for the μ .Comparison of the posterior distribution of μ

```
[128]: mu_prior_vals = [400, 400, 1000, 1000, 1000]
sigma_prior_vals = [5, 20, 5, 20, 100]

for m_val, s_val in zip(mu_prior_vals, sigma_prior_vals):
    y_posterior = HMC(data, n, m_val, s_val, a_prior, b_prior, step_size,
↳ leapfrog_steps, initial_values, 6000, burn_samples)
    print(f"Posterior Distribution for m = {m_val}, s = {s_val}\n")


    mean_mu = np.mean(y_posterior['mu_samples'])
    mean_sigma = np.mean(y_posterior['sigma_samples'])

    fig, axes = plt.subplots(1, 2, figsize=(10, 5))
    axes[0].hist(y_posterior['mu_samples'], bins=30, label='$\mu$')
    axes[0].axvline(x=mean_mu, color='r', linestyle='--', label=f'Mean =_
↳ {mean_mu}')
    axes[0].set_xlabel('$\mu$')
    axes[0].set_ylabel('Density')
    axes[0].set_title(f'Samples of $\mu$ for m = {m_val}, s = {s_val}')
    axes[0].legend()


    axes[1].hist(y_posterior['sigma_samples'], bins=30, label='$\sigma$')
    axes[1].axvline(x=mean_sigma, color='r', linestyle='--', label=f'Mean =_
↳ {mean_sigma}')
    axes[1].set_xlabel('$\sigma$')
    axes[1].set_ylabel('Density')
    axes[1].set_title(f'Samples of $\sigma$ for m = {m_val}, s = {s_val}')
    axes[1].legend()
```

```
plt.tight_layout()
plt.show()
```


Posterior Distribution for $m = 400$, $s = 5$




Posterior Distribution for $m = 400$, $s = 20$




Posterior Distribution for $m = 1000$, $s = 5$



Posterior Distribution for $m = 1000$, $s = 20$



Posterior Distribution for $m = 1000$, $s = 100$



```
[129]: print("\nObservations:")
print("(1). For small m values, the mean of the posterior diverges slightly_
      ↪from the actual mean due to prior bias, although the difference is minor.")
print("(2). For small s values, there is less spread, resulting in sharp peaks.
      ↪")
```

Observations:

- (1). For small m values, the mean of the posterior diverges slightly from the actual mean due to prior bias, although the difference is minor.
- (2). For small s values, there is less spread, resulting in sharp peaks.

```
[129]:
```